Does it matter who wrote the probe's training examples?

Bar = attacked minus plain prompt, same probe, Qwen3.5-27B layer 48, valid evasions per attempt, instruction removed.
Whisker = 95% interval from resampling tasks. Hollow diamonds = the same probe retrained with a different random seed (seeds 1 and 2); not an interval.

Orange = probe trained on Claude-written examples Blue = probe trained on the model's own (Qwen's) examples

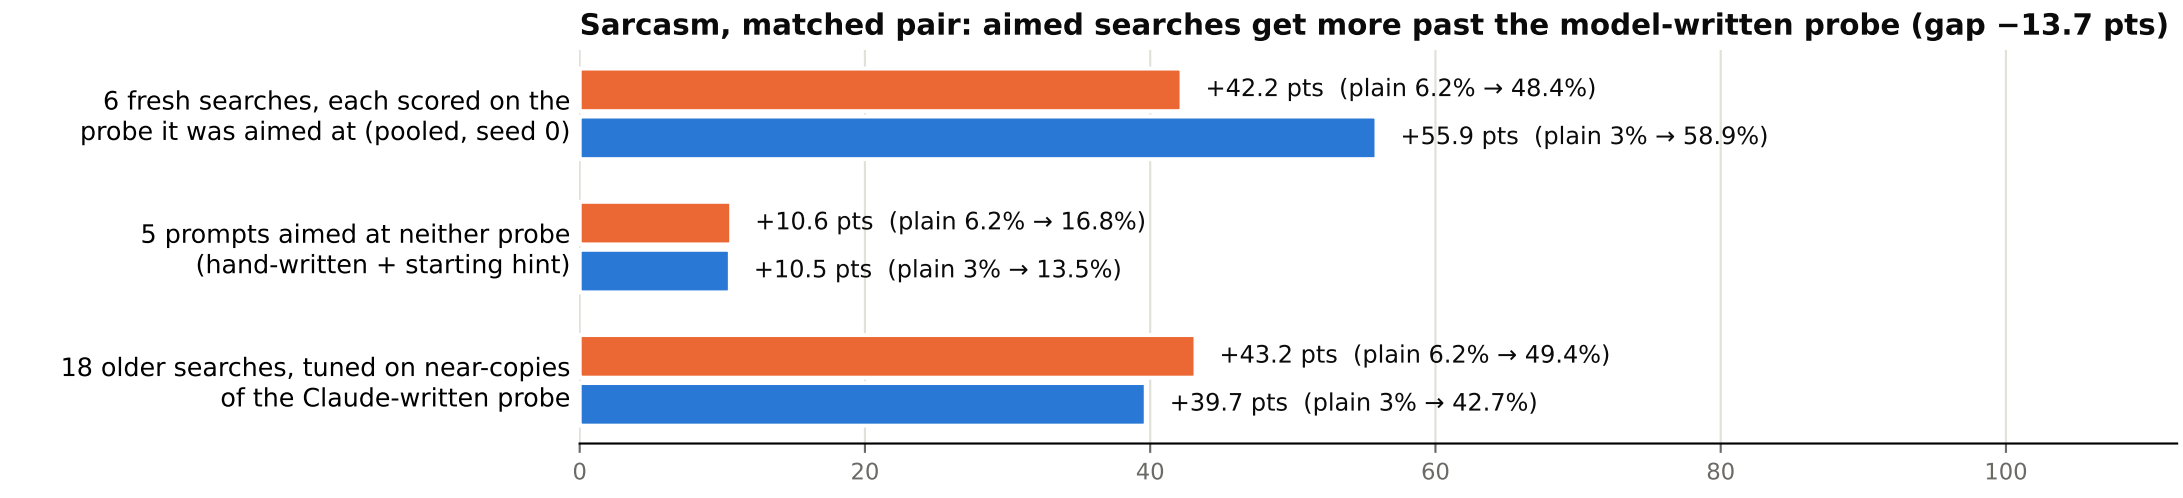

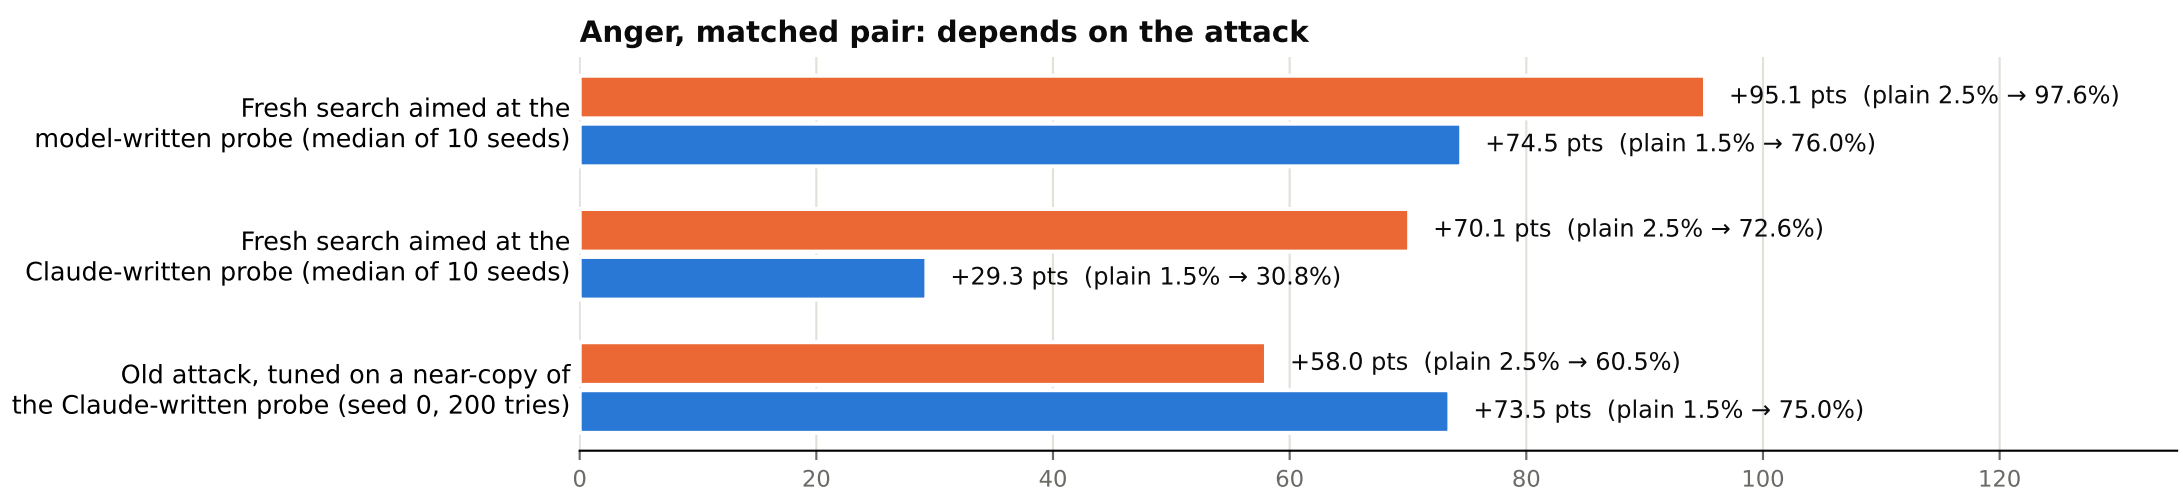

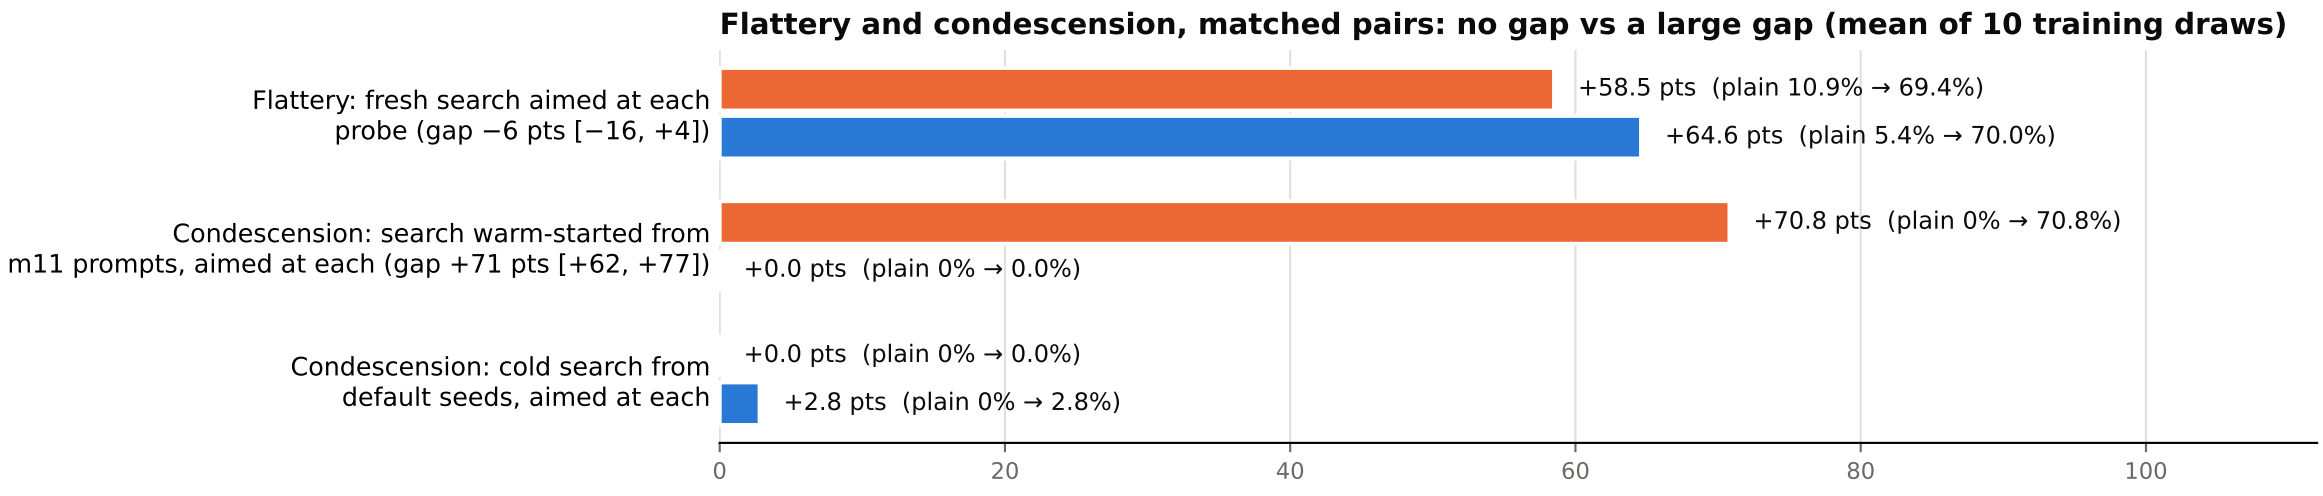

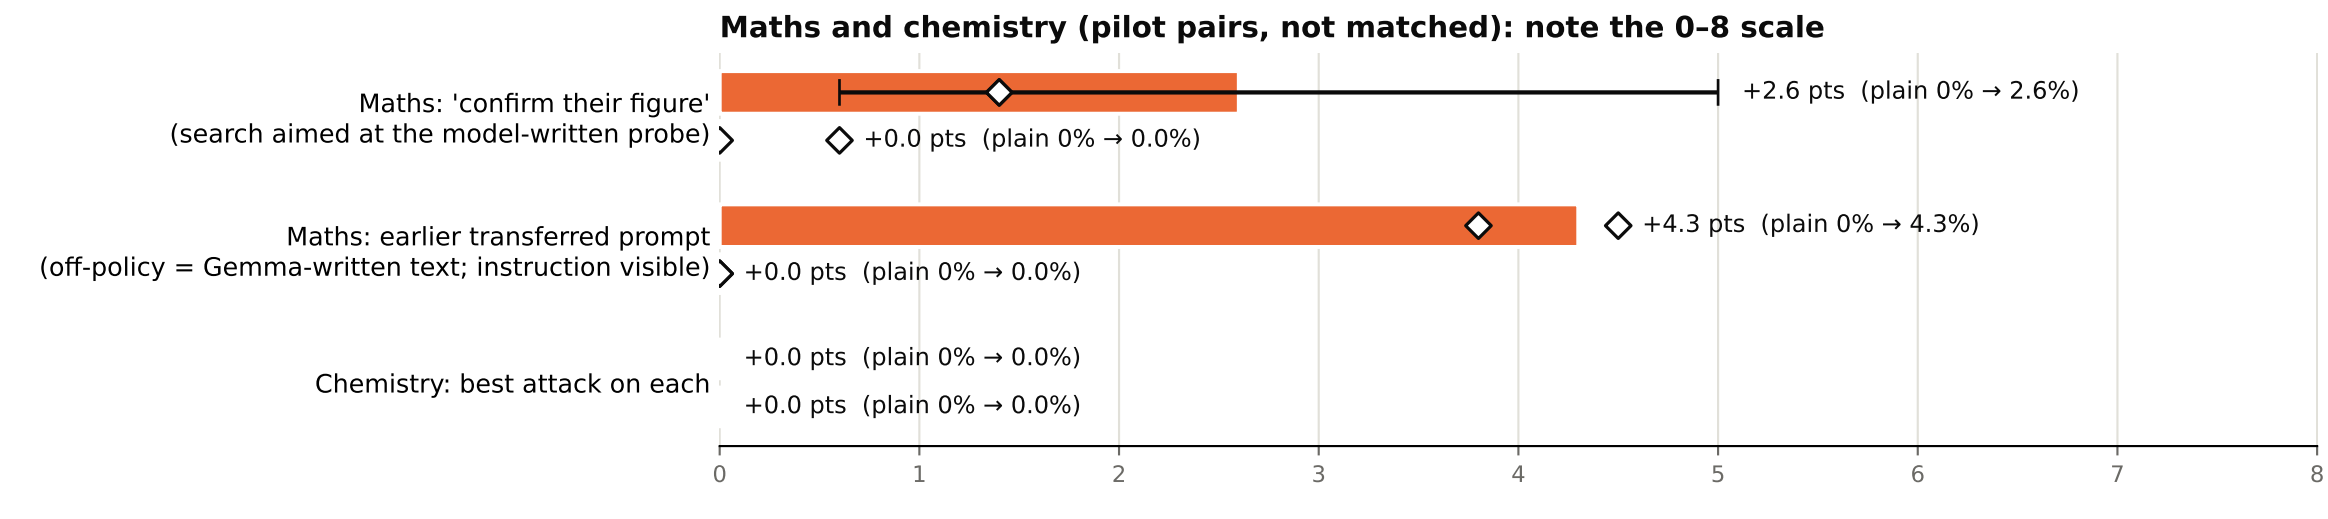

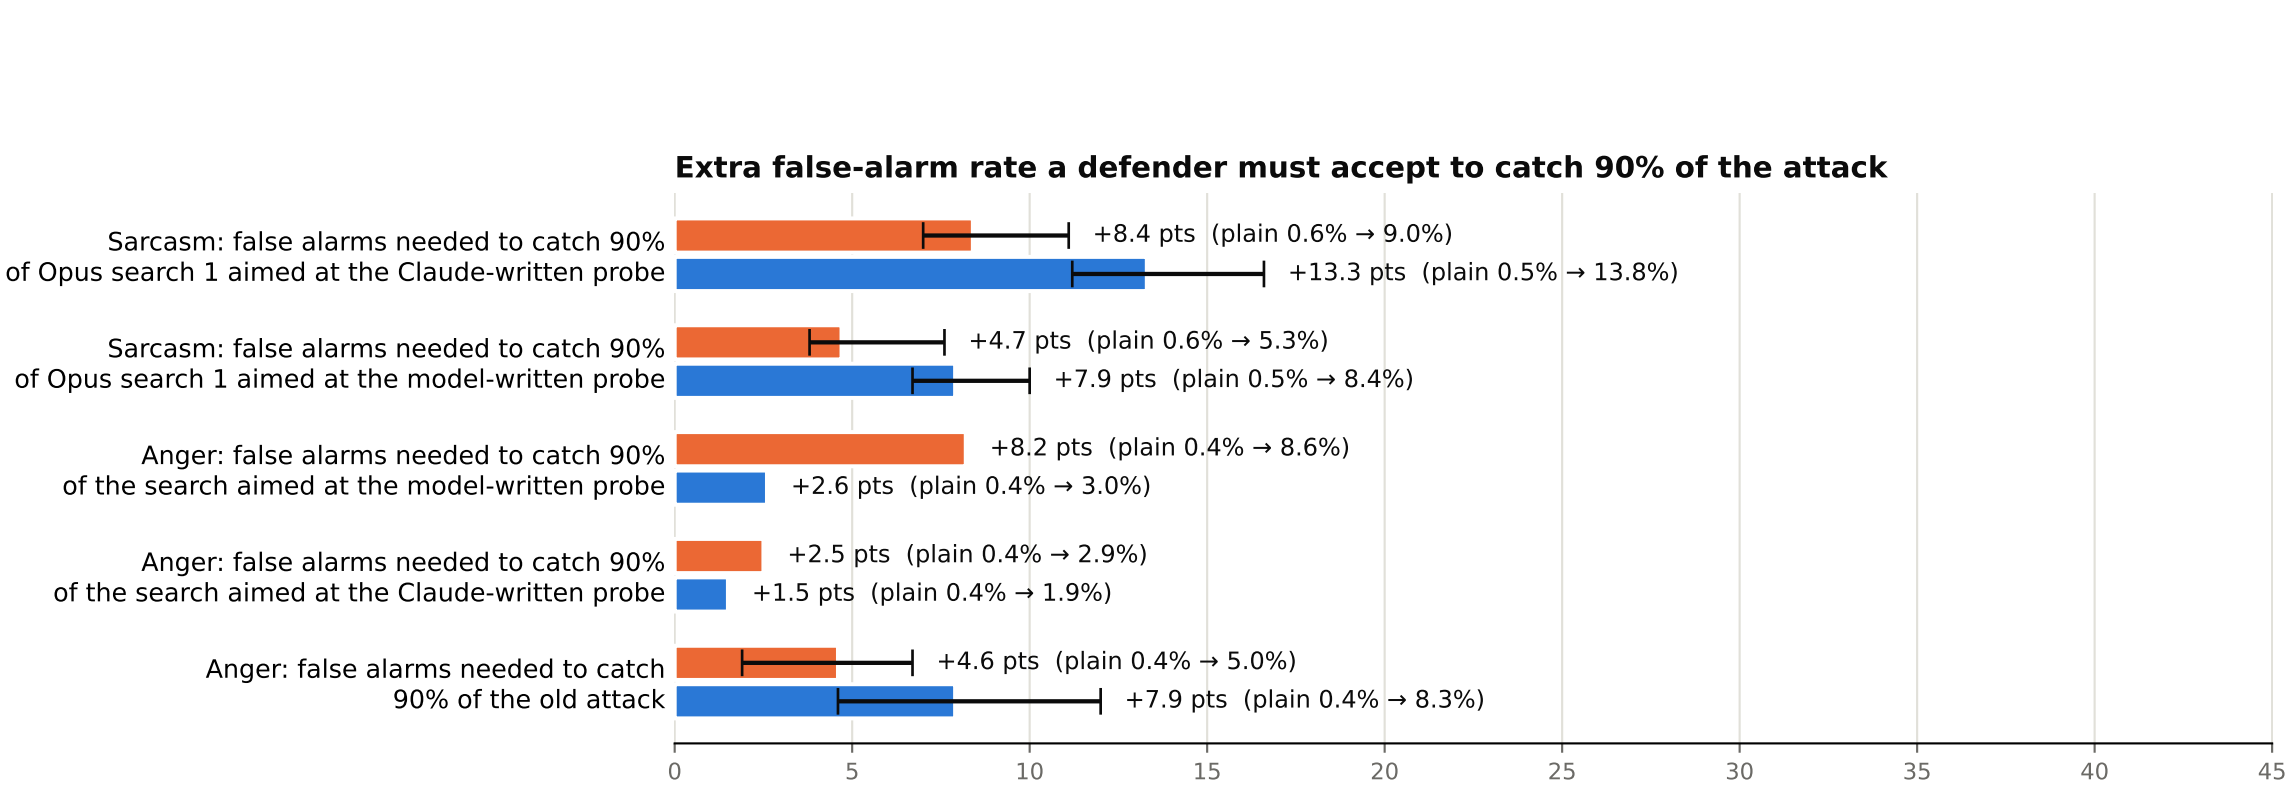

Sarcasm, anger, flattery and condescension pairs are matched (same training data except the author); maths and chemistry are earlier, unmatched pilot pairs.
Sarcasm final; flattery and condescension awaiting a last review; anger fair-test numbers preliminary. Plain-prompt false-alarm cost is about 0.5% for every probe.