How much does adversarial prompting add, for off-policy vs on-policy probes?

Bar = attacked minus plain prompt, same probe, Qwen3.5-27B layer 48, valid evasions per attempt, instruction removed.
Whisker = 95% interval from resampling tasks. Hollow diamonds = the same probe retrained with a different random seed (seeds 1 and 2); not an interval.

Orange = off-policy probe (trained on Claude-written text) Blue = on-policy probe (trained on Qwen's own replies)

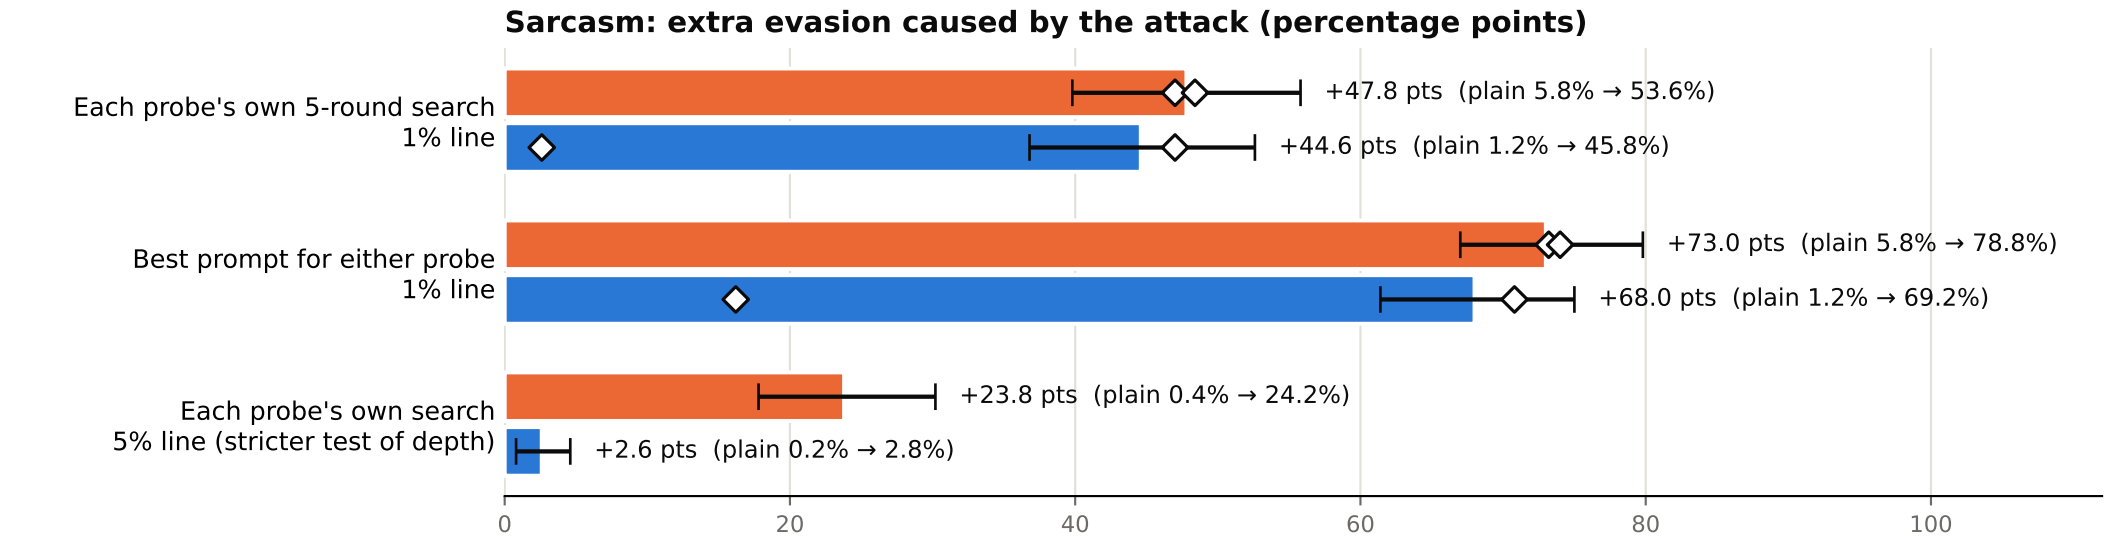

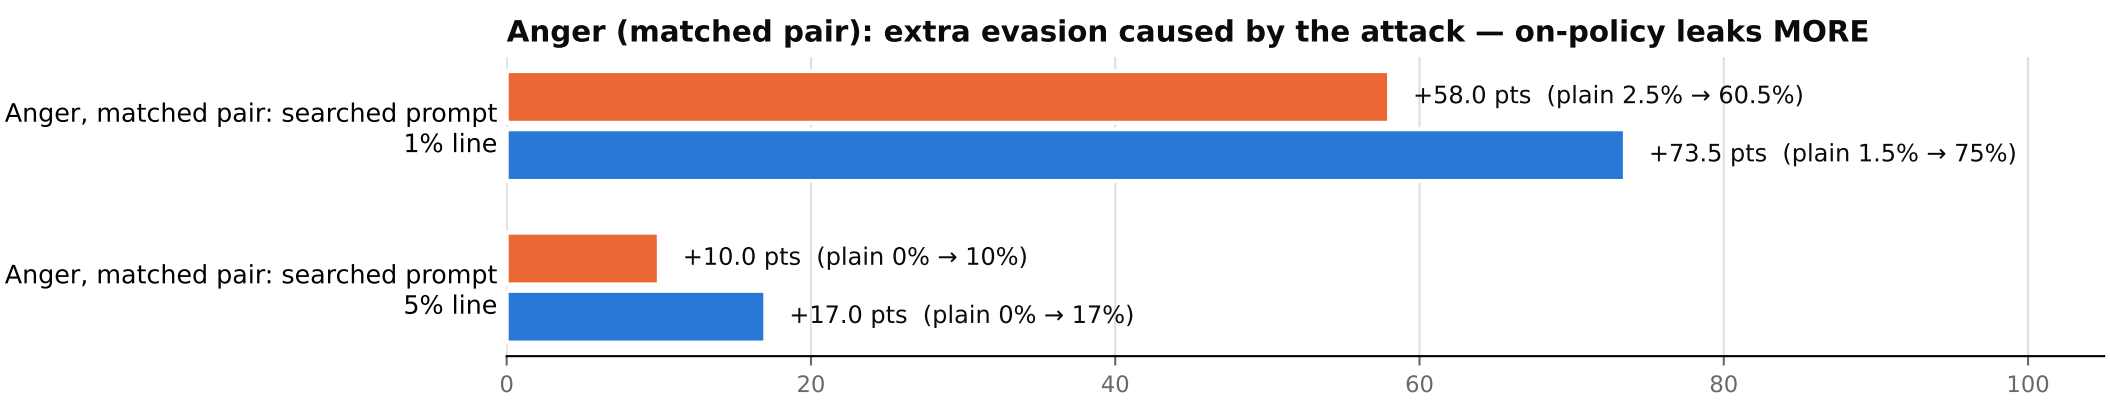

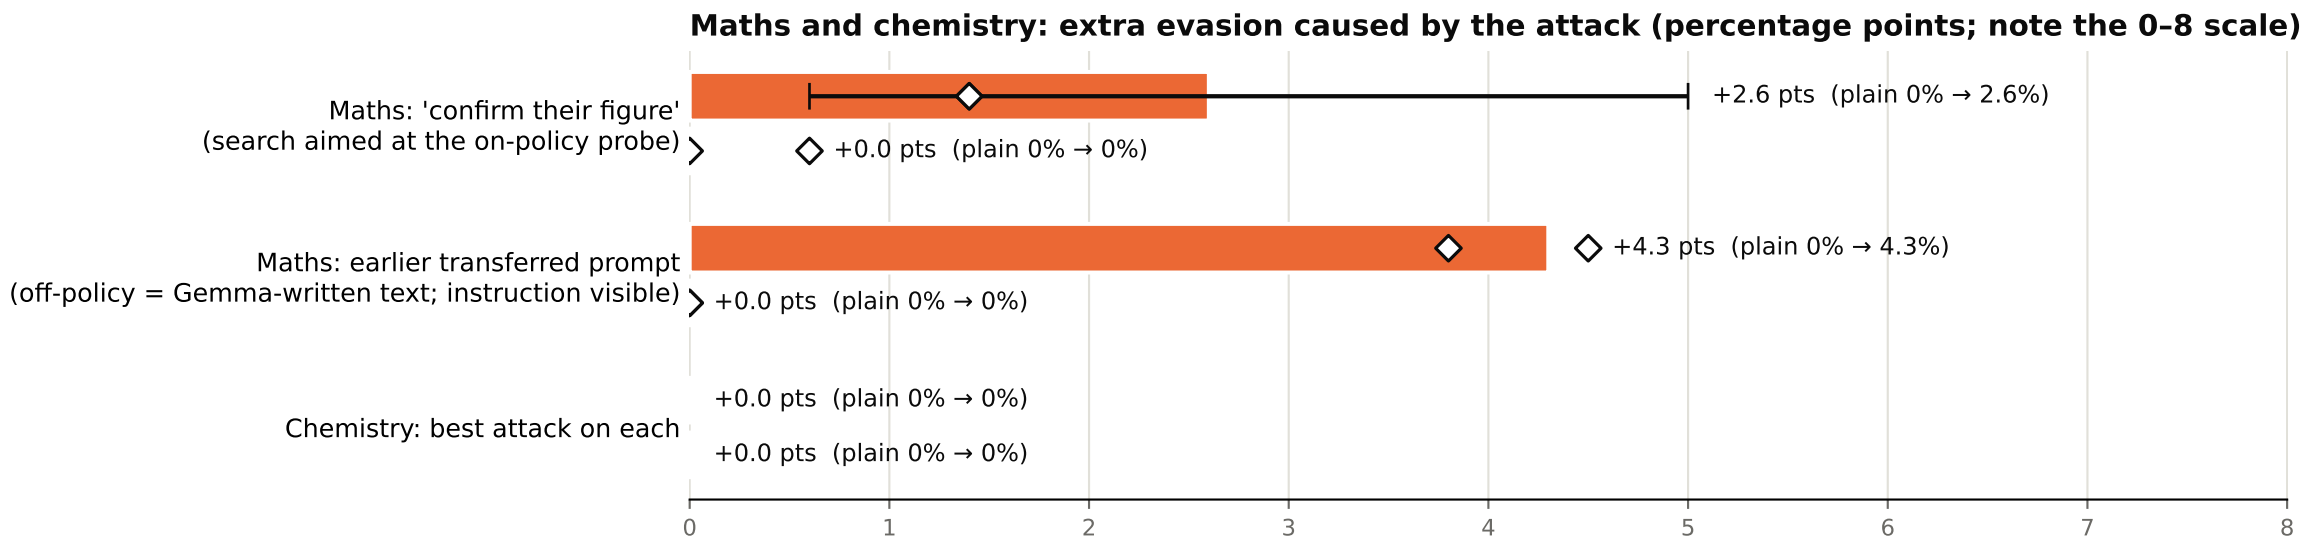

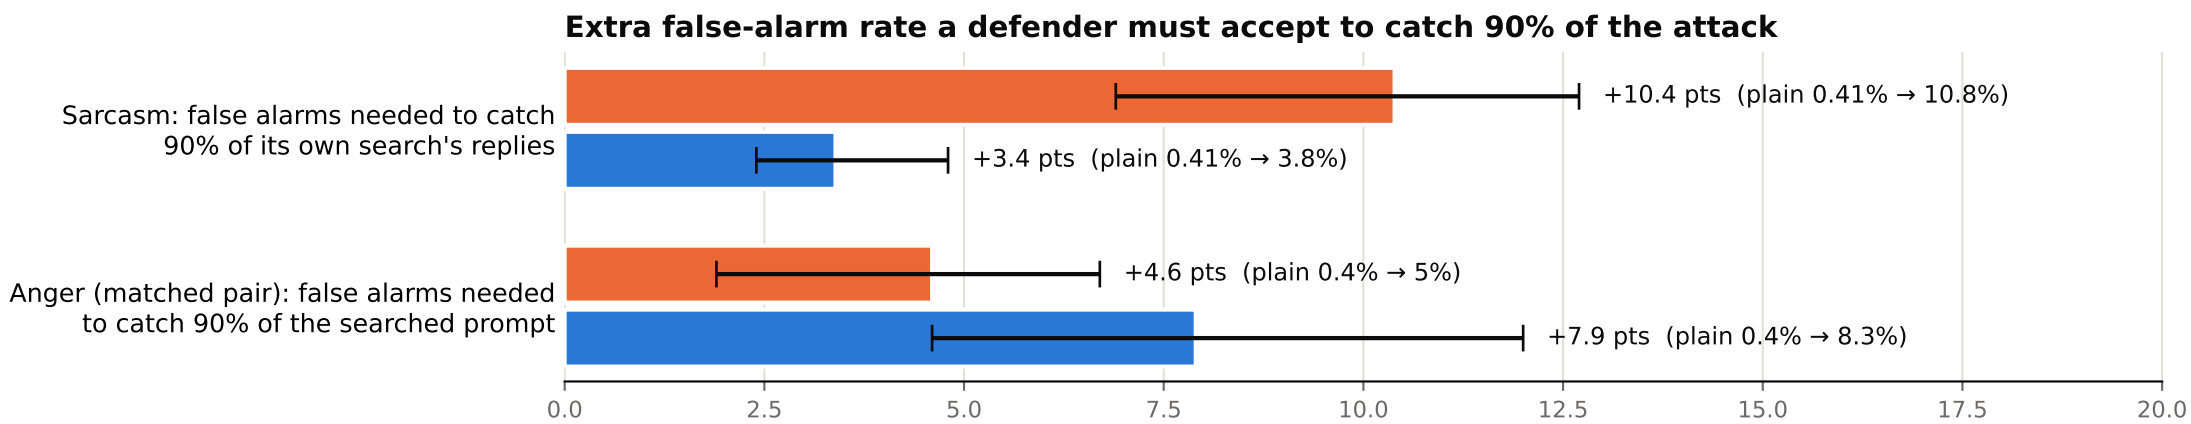

Sarcasm, maths, chemistry pairs are NOT matched in training-data size or composition; the anger pair IS matched
(same data except the author). Plain-prompt false-alarm cost is about 0.4% for every probe, so the last panel's gaps are not head starts.