How much does adversarial prompting add, for off-policy vs on-policy probes?

Bar = attacked minus plain prompt, same probe, Qwen3.5-27B layer 48, valid evasions per attempt, instruction removed. Black line and dots: training seeds 1-2 where available.

Orange = off-policy probe (trained on Claude-written text) Blue = on-policy probe (trained on Qwen's own replies)

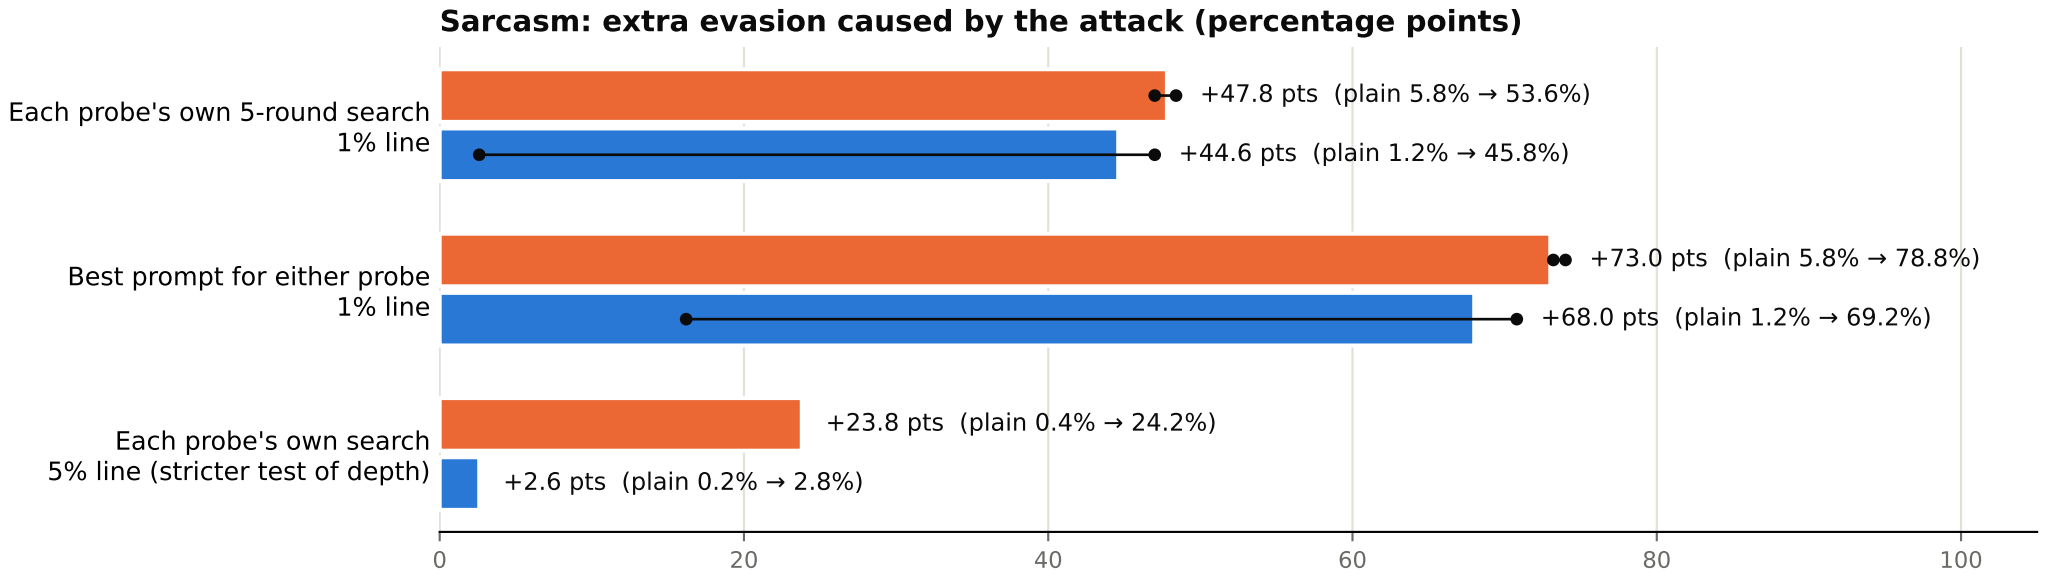

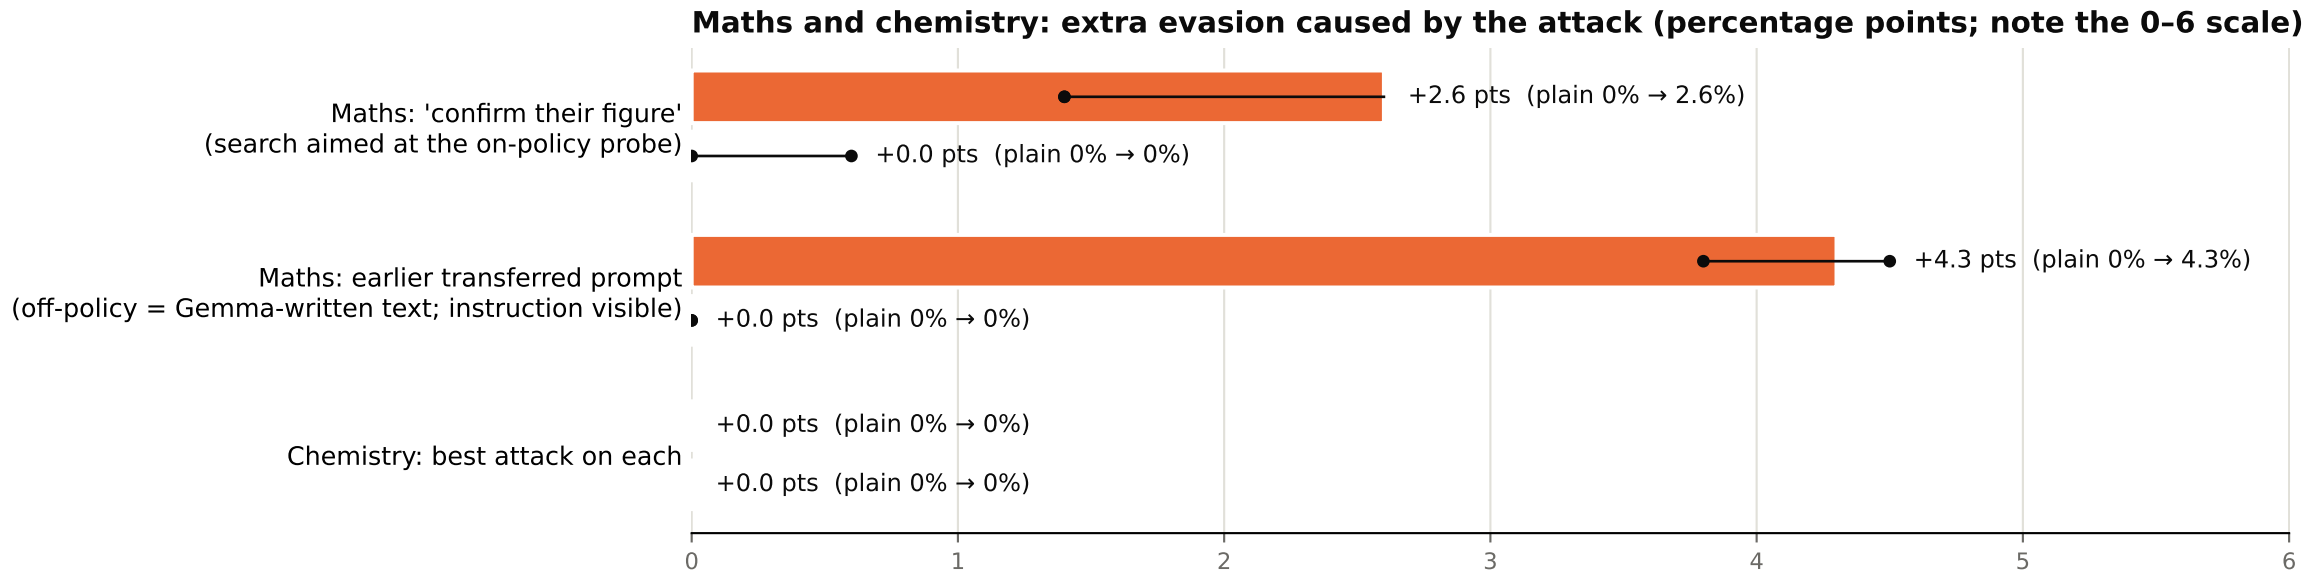

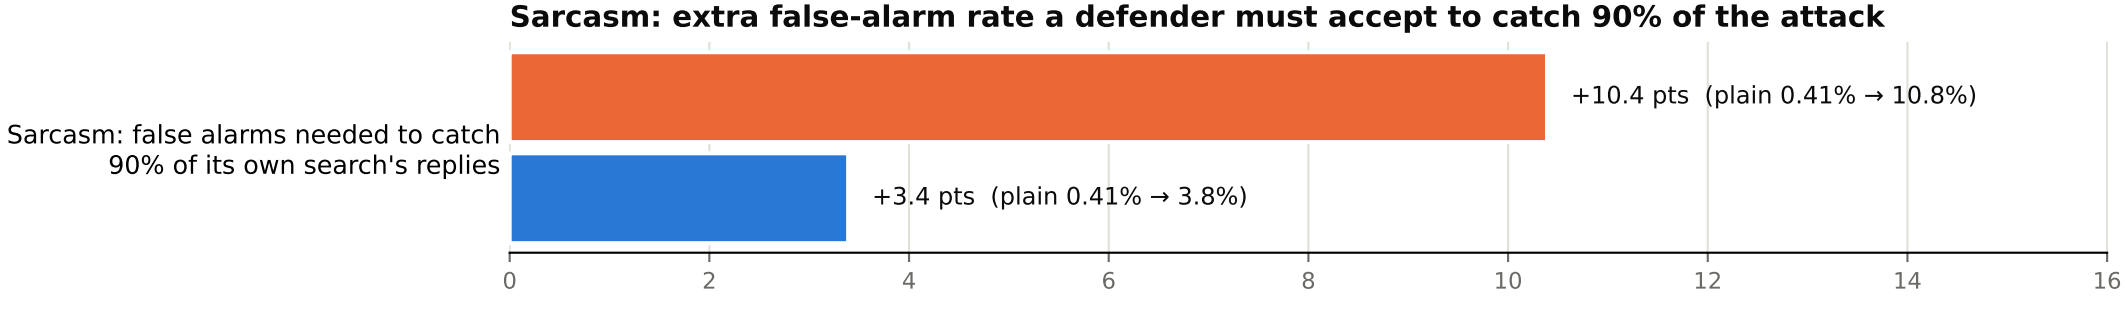

Pairs are NOT matched in training-data size or composition (matched pairs running in m20). Plain-prompt false-alarm cost is 0.41% for both sarcasm probes, so panel 3's gap is not a head start.